

The screenshot shows a GitHub repository page for 'BLENDED_LEARNING_Implementation-of-Linear-and-Polynomial-Regression-Models-for-Predicting-Car-Prices'. The repository has 98 lines (77 loc) and 3.26 KB. A commit by 'anisekinsella' was made at 4373fbb · now.

BLENDED_LEARNING

Implementation-of-Linear-and-Polynomial-Regression-Models-for-Predicting-Car-Prices

AIM:

To write a program to predict car prices using Linear Regression and Polynomial Regression models.

Equipments Required:

1. Hardware – PCs
2. Anaconda – Python 3.7 Installation / Jupyter notebook

Algorithm

1. Load the car price dataset, preprocess the data (handle missing values, encode categorical variables), and split it into training and testing sets.
2. Train a Linear Regression model using the training data and predict car prices for the test data.
3. Transform features into polynomial features, train a Polynomial Regression model, and predict prices.
4. Evaluate and compare both models using metrics like MAE, MSE, and R² score to select the best model.

Program:

```
import pandas as pd
from sklearn.model_selection import train_test_split
```

```

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import mean_absolute_error, r2_score
import matplotlib.pyplot as plt

# Load data
df = pd.read_csv('encoded_car_data (1).csv')
print(df.head())

# Select features & target
X = df.drop('price', axis=1)
y = df['price']
print(df.head())

# Select features & target
X = df[['enginesize', 'horsepower', 'citympg', 'highwaympg']]
y = df['price']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 1. Linear Regression (with scaling)
lr = Pipeline([
    ('scaler', StandardScaler()),
    ('model', LinearRegression())
])
lr.fit(X_train, y_train)
y_pred_linear = lr.predict(X_test)

# 2. Polynomial Regression (degree=2)
poly_model = Pipeline([
    ('poly', PolynomialFeatures(degree=2)),
    ('scaler', StandardScaler()),
    ('model', LinearRegression())
])
poly_model.fit(X_train, y_train)
y_pred_poly = poly_model.predict(X_test)

# Evaluate models
print('Name: Anise Kinsella A')
print('Reg. No: 212225040021')
print("Linear Regression:")
mse=mean_squared_error(y_test,y_pred_linear)
print('MSE= ',mean_squared_error(y_test,y_pred_linear))
r2score=r2_score(y_test,y_pred_linear)
print('MAE= ',mean_absolute_error(y_test,y_pred_linear))
r2score=r2_score(y_test,y_pred_linear)

print("\nPolynomial Regression:")
print(f"MSE: {mean_squared_error(y_test, y_pred_poly):.2f}")
print(f"R^2: {r2_score(y_test, y_pred_poly):.2f}")

# Plot actual vs predicted
plt.figure(figsize=(10, 5))
plt.scatter(y_test, y_pred_linear, label='Linear', alpha=0.6)
plt.scatter(y_test, y_pred_poly, label='Polynomial (degree=2)', alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', label='Perfect Prediction')

```

```
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Linear vs Polynomial Regression")
plt.legend()
plt.show()
```

Output:

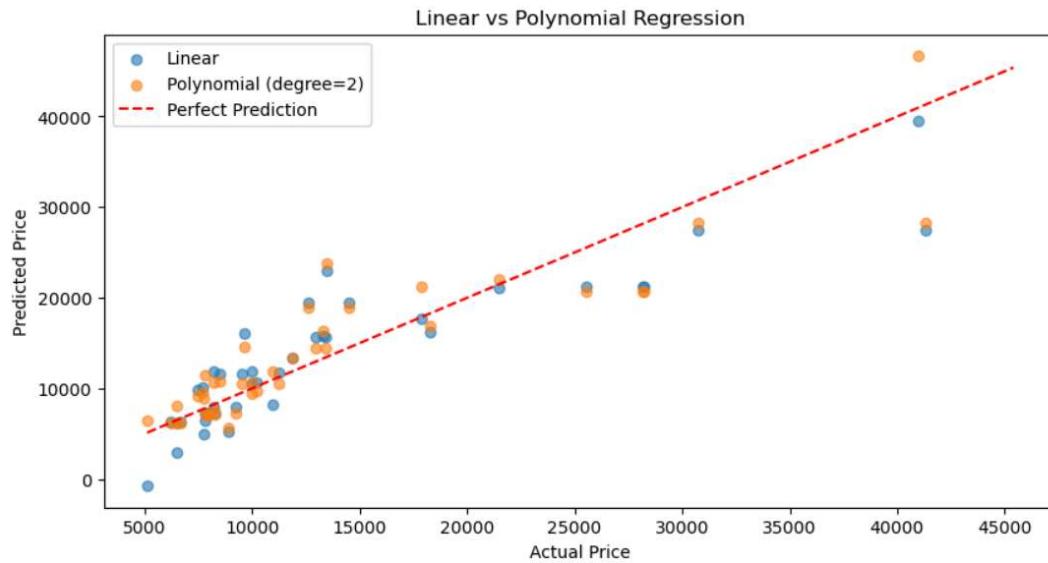
```
diesel   gas   std   turbo   convertible   hardtop   hatchback   sedan   wagon   \
0       0.0   1.0   1.0     0.0           1.0       0.0         0.0     0.0     0.0
1       0.0   1.0   1.0     0.0           1.0       0.0         0.0     0.0     0.0
2       0.0   1.0   1.0     0.0           0.0       0.0         0.0     1.0     0.0
3       0.0   1.0   1.0     0.0           0.0       0.0         0.0     0.0     1.0
4       0.0   1.0   1.0     0.0           0.0       0.0         0.0     0.0     1.0

4wd    ...   wheelbase   curbweight   enginesize   boreratio   horsepower   \
0   0.0   ...        88.6      2548.0      130.0       3.47      111.0
1   0.0   ...        88.6      2548.0      130.0       3.47      111.0
2   0.0   ...        94.5      2823.0      152.0       2.68      154.0
3   0.0   ...        99.8      2337.0      109.0       3.19      102.0
4   1.0   ...        99.4      2824.0      136.0       3.19      115.0

carlength   carwidth   citympg   highwaympg   price
0       168.8      64.1       21.0        27.0    13495.0
1       168.8      64.1       21.0        27.0    16500.0
2       171.2      65.5       19.0        26.0    16500.0
3       176.6      66.2       24.0        30.0    13950.0
4       176.6      66.4       18.0        22.0    17450.0

[5 rows x 36 columns]
Name: Anise Kinsella A
Reg. No: 212225040021
Linear Regression:
MSE= 16471505.900042146
MAE= 2892.628134137953

Polynomial Regression:
MSE: 15247661.89
R2: 0.81
```



Result:

Thus, the program to implement Linear and Polynomial Regression models for predicting car prices was written and verified using Python programming.

BLENDED_LEARNING_Implementation-of-Linear-and-Polynomial-Regression-Models-for-Predicting-Car-Prices / README.md ↑ Top

Preview Code Blame

Raw

This screenshot shows a GitHub repository page for a README.md file. At the top, there's a breadcrumb navigation: 'BLENDED_LEARNING_Implementation-of-Linear-and-Polynomial-Regression-Models-for-Predicting-Car-Prices' followed by '/ README.md'. To the right of the breadcrumb is a 'Top' link with an upward arrow icon. Below the breadcrumb, there are three buttons: 'Preview', 'Code', and 'Blame'. To the right of these buttons is a horizontal row of icons: a folder icon with a double-headed arrow, a copy icon, a raw file icon, a download icon, an edit icon, and a dropdown menu icon. At the bottom of the screenshot, there's a small portion of the page content which is mostly blank.